

File System Calls – Part 2

umask

- `#include <sys/types.h>`
`#include <sys/stat.h>`
`mode_t umask(mode_t mask);`
 - Set file mode creation *mask* and return the old value.
 - When creating a file, permissions are turned off if the corresponding bits in *mask* are set.
 - Return value
 - This system call always succeeds and the previous value of the mask is returned.
 - cf. “umask” shell command

Example #5: umask

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <sys/types.h>
4 #include <sys/stat.h>
5 #include <fcntl.h>
6
7 int main(void) {
8     umask(0);
9     if (creat("foo", S_IRUSR|S_IWUSR|S_IRGRP|S_IWGRP|S_IROTH|S_IWOTH) < 0) {
10         perror("creat error for foo");
11         exit(1);
12     }
13
14     umask(S_IRGRP|S_IWGRP|S_IROTH|S_IWOTH);
15     if (creat("bar", S_IRUSR|S_IWUSR|S_IRGRP|S_IWGRP|S_IROTH|S_IWOTH) < 0) {
16         perror("creat error for bar");
17         exit(1);
18     }
19     exit(0);
20 }
```

```
> ./main
> ls -la foo
-rw-rw-rw- 1 runner runner 0 Mar 30 00:24 foo
> ls -la bar
-rw----- 1 runner runner 0 Mar 30 00:24 bar
```

Example #5: umask



```
main.c saving...
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <sys/types.h>
4 #include <sys/stat.h>
5 #include <fcntl.h>
6
7 int main(void){
8     umask(0);
9     if (creat("foo", S_IRUSR|S_IWUSR|S_IRGRP|S_IWGRP|S_IROTH|
10         S_IWOTH) < 0) {
11         perror("creat error for foo");
12         exit(1);
13     }
14
15     umask(S_IRGRP|S_IWGRP|S_IROTH|S_IWOTH);
16     if (creat("bar", S_IRUSR|S_IWUSR|S_IRGRP|S_IWGRP|S_IROTH|
17         S_IWOTH) < 0) {
18         perror("creat error for bar");
19         exit(1);
20     }
21     exit(0);
22 }
```

```
https://yeonseub.hasooeun.repl.run
clang version 7.0.0-3-ubuntu0.18.04.1 (tags/RELEASE_700/final)
> clang-7 -pthread -lm -o main main.c
> ./main
> []
```

access (1)

- `#include <unistd.h>`

`int access(const char *pathname, int mode);`

- Checks whether the process would be allowed to read, write or test for existence of the file whose name is *pathname*.
- If *pathname* is a symbolic link, permissions of the file referred to by this symbolic link are tested.
- The check is done with the process's real uid and gid. (neither effective uid nor effective gid)

access (2)

- *mode*
 - R_OK
 - file exist, read permission
 - W_OK
 - file exist, write permission
 - X_OK
 - file exist, execution permission
 - F_OK
 - file exist
- Return value
 - On success zero is returned. On error -1 is returned.
 - Returns an error if any of the access types in the requested call fails, even if other types might be successful.
 - Ex: access("bit", R_OK|W_OK) checks read and write permission.

Example #6: access

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <sys/types.h>
5 #include <fcntl.h>
6
7 int main(int argc, char *argv[]) {
8     if (argc != 2) {
9         fprintf(stderr, "usage: a. out <pathname>");
10        exit(1);
11    }
12    if (access(argv[1], R_OK) < 0)
13        perror("access error");
14    else
15        printf("read access OK\n");
16    if (open(argv[1], O_RDONLY) < 0)
17        perror("open error");
18    else
19        printf("open for reading OK\n");
20    exit(0);
21 }
```

```
➤ ./main
usage: ./main <pathname>exit status 1
```

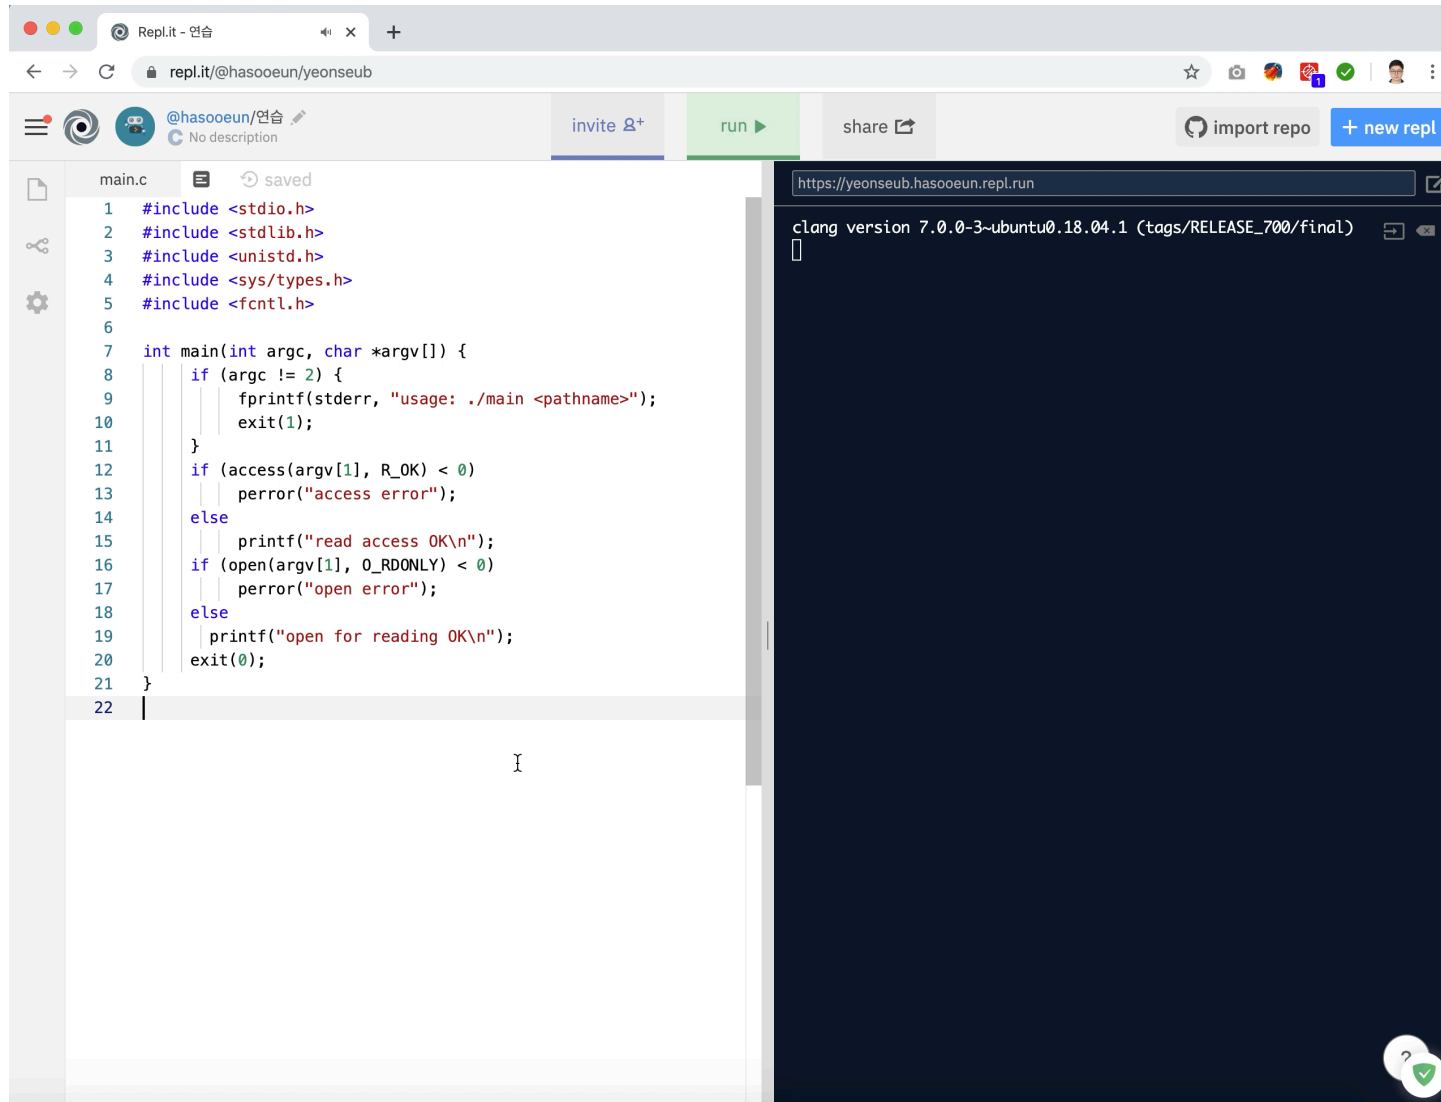
```
➤ ls -la foo
-rw-rw-rw- 1 runner runner 0 Mar 30 00:24 foo
```

```
➤ ./main foo
read access OK
open for reading OK
```

```
➤ chmod 0266 foo
➤ ls -la foo
--w-rw-rw- 1 runner runner 0 Mar 30 00:24 foo
```

```
➤ ./main foo
access error: Permission denied
open error: Permission denied
```

Example #6: access



The screenshot shows a Repl.it IDE interface. The left pane displays a C program named `main.c` with the following code:

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <sys/types.h>
5 #include <fcntl.h>
6
7 int main(int argc, char *argv[]) {
8     if (argc != 2) {
9         fprintf(stderr, "usage: ./main <pathname>");
10        exit(1);
11    }
12    if (access(argv[1], R_OK) < 0)
13        perror("access error");
14    else
15        printf("read access OK\n");
16    if (open(argv[1], O_RDONLY) < 0)
17        perror("open error");
18    else
19        printf("open for reading OK\n");
20    exit(0);
21 }
22
```

The right pane shows the terminal output, which includes the clang version information and a prompt for input:

```
clang version 7.0.0-3-ubuntu0.18.04.1 (tags/RELEASE_700/final)
[ ]
```


stat, fstat, lstat

- `#include <sys/stat.h>`
`#include <unistd.h>`
`int stat(const char *file_name, struct stat *buf);`
`int fstat(int fd, struct stat *buf);`
`int lstat(const char *file_name, struct stat *buf);`
 - Return information about the specified file.
 - Do not need any access rights to the file. But need search rights to all directories named in the path leading to the file.
 - `stat()`, `fstat()`
 - Stats the file pointed to by `file_name` or by `fd` and fills in `buf`.
 - `lstat()`
 - Same as `stat()` except that the symbolic link is stated itself (i.e. do not follow the link).

Structure stat

```
struct stat {  
    dev_t st_dev;           /* device number */  
    ino_t st_ino;           /* inode number */  
    mode_t st_mode;         /* file type, mode (permissions) */  
    nlink_t st_nlink;       /* number of hard links */  
    uid_t st_uid;           /* user ID of owner */  
    gid_t st_gid;           /* group ID of owner */  
    dev_t st_rdev;          /* device type for special files (if inode device) */  
    off_t st_size;          /* total size, in bytes */  
    unsigned long st_blksize; /* blocksize for I/O */  
    unsigned long st_blocks; /* # of blocks allocated */  
    time_t st_atime;         /* time of last access */  
    time_t st_mtime;         /* time of last modification */  
    time_t st_ctime;         /* time of last status change */  
};
```

stat Fields

- `st_blocks`
 - The number of blocks allocated to this file.
- `st_blksize`
 - The “preferred” block size for efficient file system I/O.
 - Writing to a file in smaller chunks may cause an inefficient *read-modify-rewrite*.
- `st_atime`
 - Changed by `mknod()`, `utime()`, `read()`, `write()`, and `truncate()`.
- `st_mtime`
 - Changed by `mknod()`, `utime()`, and `write()`.
- `st_ctime`
 - Changed by writing or by setting inode information.
 - owner, group, link count, mode, etc.

POSIX Macro for File Type

- `S_ISLNK(st_mode)`: symbolic link
- `S_ISREG(st_mode)`: regular file
- `S_ISDIR(st_mode)`: directory
- `S_ISCHR(st_mode)`: character device
- `S_ISBLK(st_mode)`: block device
- `S_ISFIFO(st_mode)`: fifo (named pipe)
- `S_ISSOCK(st_mode)`: socket (unix domain socket)

POSIX Flags for st_mode Field

- S_IFMT: 0170000 bit mask for the file type bit fields.
- S_IFSOCK: 0140000 socket.
- S_IFLNK: 0120000 symbolic link.
- S_IFREG: 0100000 regular file
- S_IFBLK: 0060000 block device
- S_IFDIR: 0040000 directory
- S_IFCHR: 0020000 character device
- S_IFIFO: 0010000 fifo
- S_ISUID: 0004000 set-user-id bit (POSIX)
- S_ISGID: 0002000 set-group-id bit (POSIX)
- S_ISVTX: 0001000 sticky bit

file type (4bit)	special(3bit)	permission(9bit)
------------------	---------------	------------------

Usage: stat, fstat, lstat

- Return value
 - On success, zero is returned.
 - On error, -1 is returned.
- Example

```
...  
struct stat buf;  
  
stat("/etc/passwd", &buf);  
if (buf.st_mode & S_IRUSR) {  
...  
}
```

Example #7: lstat (1)

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  #include <sys/types.h>
5  #include <sys/stat.h>
6
7  int main(int argc, char *argv[]) {
8      int i;
9      struct stat buf;
10     char *ptr;
11     for (i = 1; i < argc; i++) {
12         printf("%s: ", argv[i]);
13         if (lstat(argv[i], &buf) < 0) {
14             perror("lstat error");
15             continue;
16         }
17     }
18     if (S_ISREG(buf.st_mode)) ptr = "regular";
19     else if (S_ISDIR(buf.st_mode)) ptr = "directory";
20     else if (S_ISCHR(buf.st_mode)) ptr = "character special";
21     else if (S_ISBLK(buf.st_mode)) ptr = "block special";
22     else if (S_ISFIFO(buf.st_mode)) ptr = "fifo";
23
24
25
```

Example #7: lstat (2)

```
26 #ifdef S_ISLNK
27     else if (S_ISLNK(buf.st_mode)) ptr = "symbolic link";
28 #endif
29 #ifdef S_ISSOCK
30     else if (S_ISSOCK(buf.st_mode)) ptr = "socket";
31 #endif
32     else ptr = "** unkown mode **";
33     printf("%s\n", ptr);
34 }
35 exit(0);
36 }
37 }
38 }
```

```
➤ ls -la | grep foo
-rwxrwsrwx 1 runner runner 0 Mar 30 00:24 foo
drwxr-xr-x 1 runner runner 0 Mar 30 02:02 foo-d
lrwxrwxrwx 1 runner runner 3 Mar 30 01:06 foo-s -> foo
➤ ./main foo foo-d foo-s
foo: regular
foo-d: directory
foo-s: symbolic link
```


chmod and fchmod

- `#include <sys/types.h>`
`#include <sys/stat.h>`
`int chmod(const char *path, mode_t mode);`
`int fchmod(int fd, mode_t mode);`
 - Change permissions of a file.
 - The mode of the file given by *path* or referenced by *fd* is changed.
 - *mode* is specified by OR'ing the following.
 - `S_ISUID`, `S_ISGID`, `S_ISVTX`, `S_I{R,W,X}{USR,GRP,OTH}`
 - On success, zero is returned. On error, -1 is returned.

Example #8: chmod (1)

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <sys/types.h>
4  #include <sys/stat.h>
5
6  int main(void) {
7      struct stat statbuf;
8
9      if (stat("foo", &statbuf) < 0) {
10         perror("stat error for foo");
11         exit(1);
12     }
13
14     /* turn on set-group-ID and turn off group-execute */
15     if (chmod("foo", (statbuf.st_mode & ~S_IXGRP) | S_ISGID) < 0) {
16         perror("chmod error for foo");
17         exit(1);
18     }
19
20 }
```

Example #8: chmod (2)

```
21  /* set absolute mode to "rw-r--r--" */
22  if (chmod("bar", S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH) < 0) {
23      perror("chmod error for bar");
24      exit(1);
25  }
26  exit(0);
27
28 }
```

```
ls -la foo bar
-rw-rw-rw- 1 runner runner 0 Mar 30 00:24 bar
-rwxrwxrwx 1 runner runner 0 Mar 30 00:24 foo
> ./main
> ls -la foo bar
-rw-r--r-- 1 runner runner 0 Mar 30 00:24 bar
-rwxrwsrwx 1 runner runner 0 Mar 30 00:24 foo
```

chown, fchown, lchown

- `#include <sys/types.h>`

`#include <unistd.h>`

`int chown(const char *path, uid_t owner, gid_t group);`

`int fchown(int fd, uid_t owner, gid_t group);`

`int lchown(const char *path, uid_t owner, gid_t group);`

- The owner of the file specified by *path* or by *fd* is changed.
- Only the superuser and the current owner may change the owner of a file.
- The owner of a file may change the group of the file to any group of which that owner is a member.
- When the owner or group of an executable file are changed by a non-superuser, the `S_ISUID` and `S_ISGID` mode bits are *cleared*.
- On success, zero is returned. On error, -1 is returned.
- *lchown()* is same as *chown()* except that the symbolic link is stated itself (i.e. do not follow the link).

truncate and ftruncate

- `#include <unistd.h>`

`int truncate(const char *path, off_t length);`

`int ftruncate(int fd, off_t length);`

- Causes the file referenced by a path or a file descriptor to have a size of *length* bytes.
- If the file previously was larger than *length*, the extra data is lost.
- If the file size was less than *length*, the effect is system dependent.
- Return value
 - On success, zero is returned.
 - On error, -1 is returned.

link



- `#include <unistd.h>`

`int link(const char *oldpath, const char *newpath);`


- Make a new name for a file
 - Creates a new link (also known as a *hard link*) to an existing file.
- If *newpath* exists it will not be overwritten.
- This new name may be used exactly as the old one for any operation.
 - Both names refer to the same file.
 - Have the same permissions and ownership.
 - It is impossible to tell which name was the original.
- Hard links, as created by `link`, cannot span file systems.

```
➤ ls -la foo-s
lrwxrwxrwx 1 runner runner 3 Mar 30 01:06 foo-s -> foo
➤ ln foo-s foo-sh
➤ ls -la | grep foo-s
lrwxrwxrwx 2 runner runner 3 Mar 30 01:06 foo-s -> foo
lrwxrwxrwx 2 runner runner 3 Mar 30 01:06 foo-sh -> foo
```



Directory link

- Only the superuser can create a hard link to a directory.



- Return value
 - On success, zero is returned.
 - On error, -1 is returned.

unlink



- `#include <unistd.h>`

`int unlink(const char *pathname);`

- Delete a name and possibly the file it refers to
 - Deletes a name from the file system.
 - If that name was the last link to a file and no processes have the file open, the file is deleted.
 - If the name was the last link to a file but any processes still have the file open, the file will remain in existence until the last file descriptor referring to it is closed.
 - If the name referred to a symbolic link, the link is removed.
- Return value
 - On success, zero is returned. On error, -1 is returned.

rename

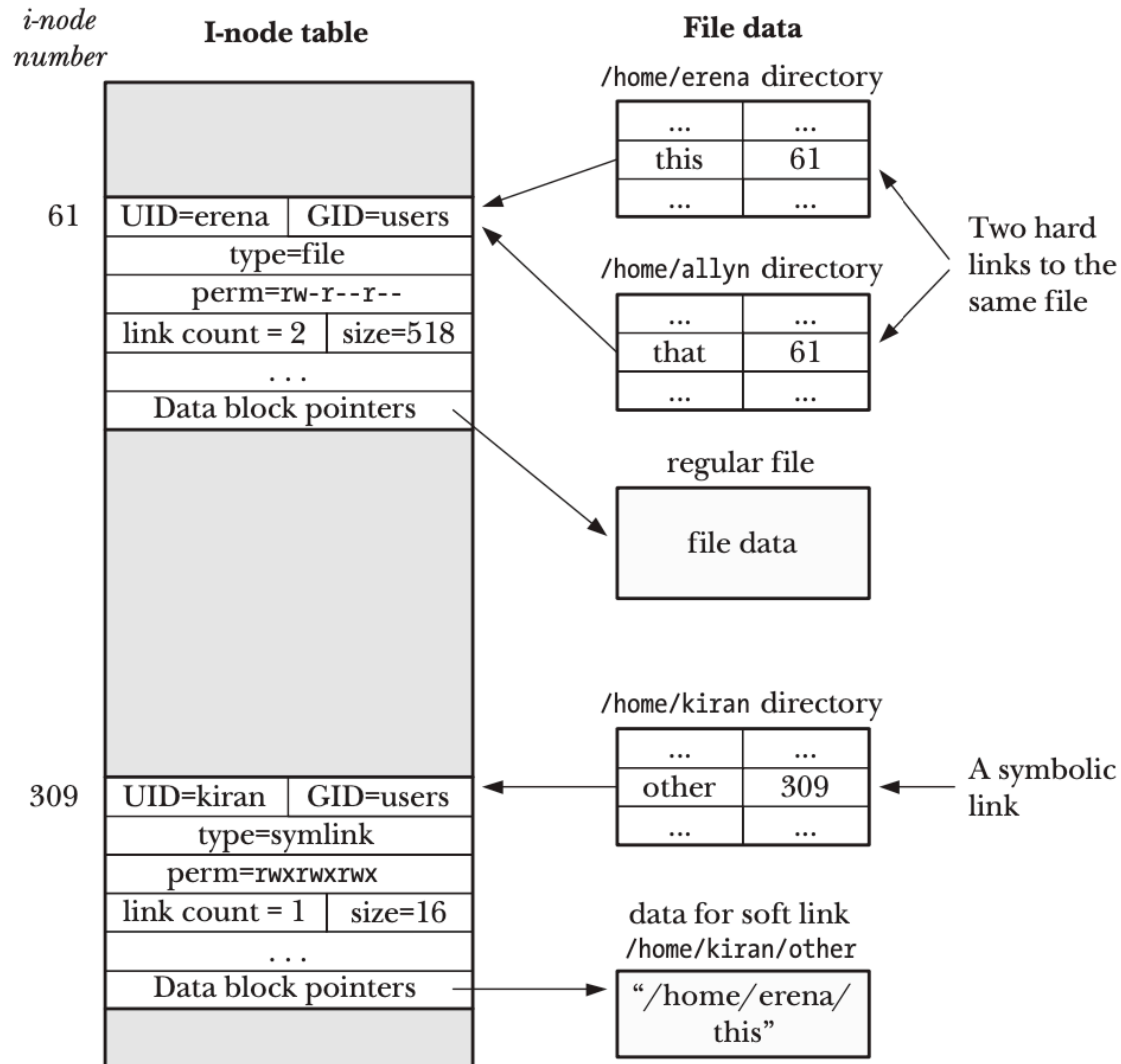
- `#include <stdio.h>`

`int rename(const char *oldpath, const char *newpath);`

- Changes the name of *oldpath* to *newpath*.
- If *oldpath* names open nonexistent file, or if *newpath* names a file that already exists, then the action of `rename()` is *implementation-dependent*.
- Return value
 - On success, zero is returned.
 - On error, -1 is returned.

- Symbolic link is indirect pointer to a file.
 - Hard link points directly to the inode of the file.
 - Limitation of hard links
 - Require that the link and the file reside in the same file system.
 - Only the superuser can create a hard link to a directory.
 - There are no file system limitations on a symbolic link.
 - Anyone can create a symbolic link to a directory.

Symbolic Link



- `#include <unistd.h>`
`int symlink(const char *oldpath, const char *newpath);`
 - Creates a symbolic link named *newpath* which contains the string *oldpath*.
 - Symbolic links are interpreted at run-time.
 - Dangling link
 - May point to a non-existing file.
 - The permissions of a symbolic link are irrelevant.
 - The ownership is ignored when following the link.
 - Permission is checked when removal or renaming of the link is requested and the link is in a directory with the sticky bit set.
 - If *newpath* exists it will not be overwritten.
 - Return value
 - On success, zero is returned. On error, -1 is returned.

- `#include <unistd.h>`

`int readlink(const char *path, char *buf, size_t bufsiz);`

- Read value of a symbolic link (does *not* follow the link).
 - Places the contents of the symbolic link *path* in the buffer *buf*, which has size *bufsiz*.
- Return value
 - The count of characters placed in the buffer if it succeeds.
 - -1 if an error occurs.

Thank you