

Process Termination (1)

- Terminating process notify its parent how it terminated.
- Normal termination
 - Pass an *exit status* as argument to `exit()` or `_exit()`.
 - Exit status is converted into a termination status by the kernel when `_exit()` is finally called.
- Abnormal termination
 - The kernel generates a termination status to indicate the reason for the abnormal termination.
- The parent of the terminated process can obtain the termination status from either `wait()` or `waitpid()`.
- *init* process becomes the parent process of any process whose parent terminates (inherited by *init*).

Process Termination (2)

- Zombie state
 - The kernel has to keep a certain amount of information consists of PID, the termination status of the process, and the amount of CPU time taken by the process.
 - The process that has terminated, but whose parent has not yet waited for it, is called a *zombie*.

exit



- `#include <stdlib.h>`

`void exit(int status);`

- Causes normal program termination and the value of *status* is returned to the parent.
 - Close and flush opened files.
 - Call functions registered in `atexit()`.
 - Call `_exit()`

atexit



- `#include <stdlib.h>`

`int atexit(void (*function)(void));`

- Registers the given *function* to be called at normal program termination.
- No arguments are passed.
- Returns the value 0 if successful; otherwise the value -1 is returned.

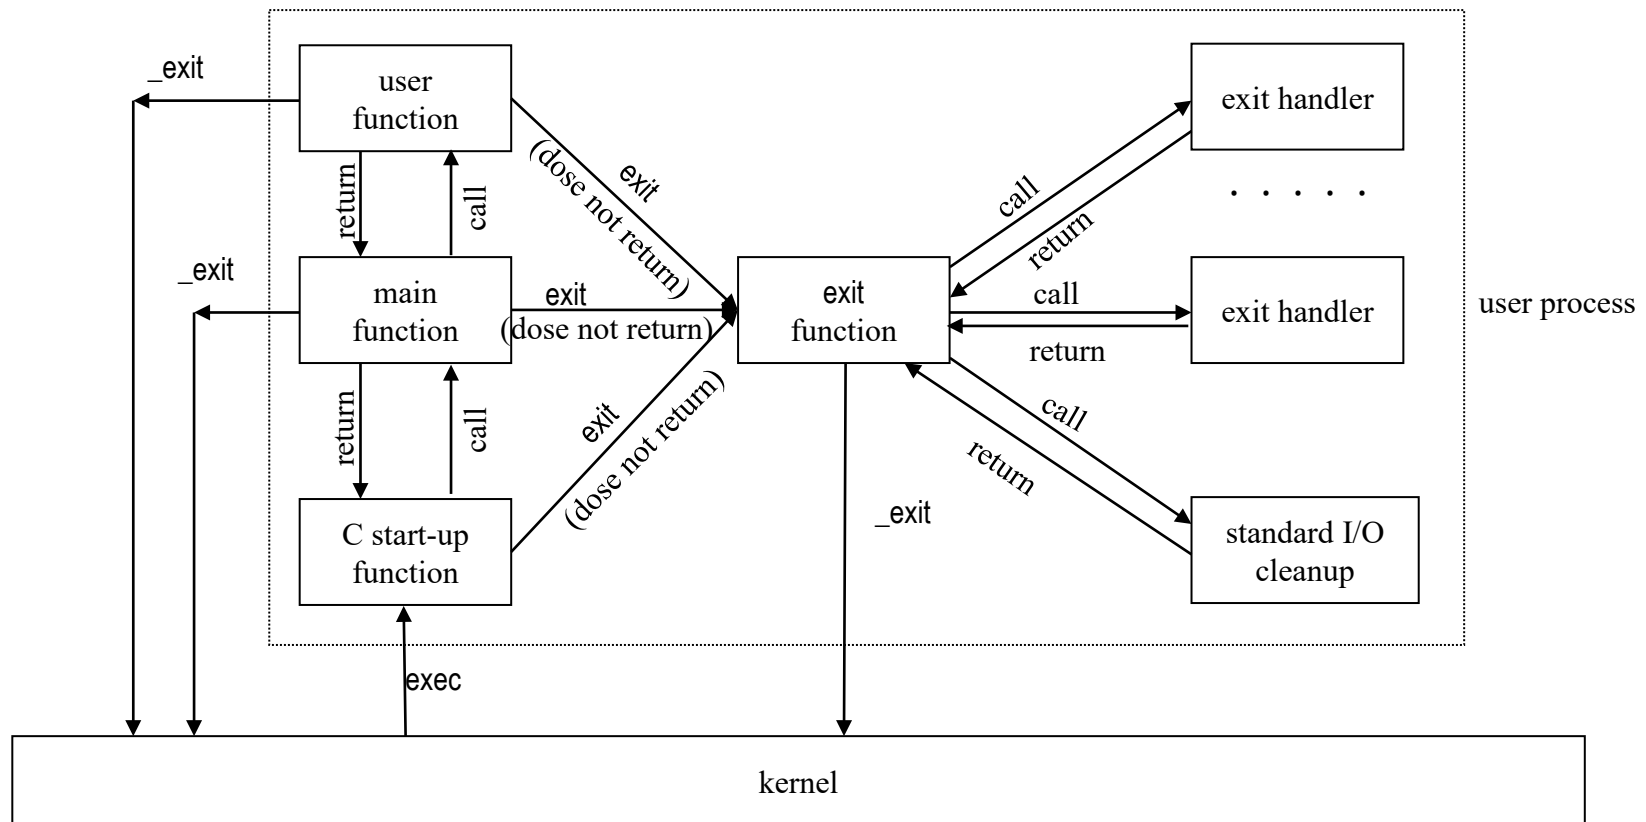
_exit

- `#include <unistd.h>`
`void _exit(int status);`
 - Terminates the calling process immediately without performing some cleanup.
 - Deallocate everything but PCB entry.
 - *status* is returned to the parent process as the process's exit status.

pointer	process state
process number(PID)	
program counter	
registers	
memory limits	
list of open files	
...	

PCB structure

Exit Functions



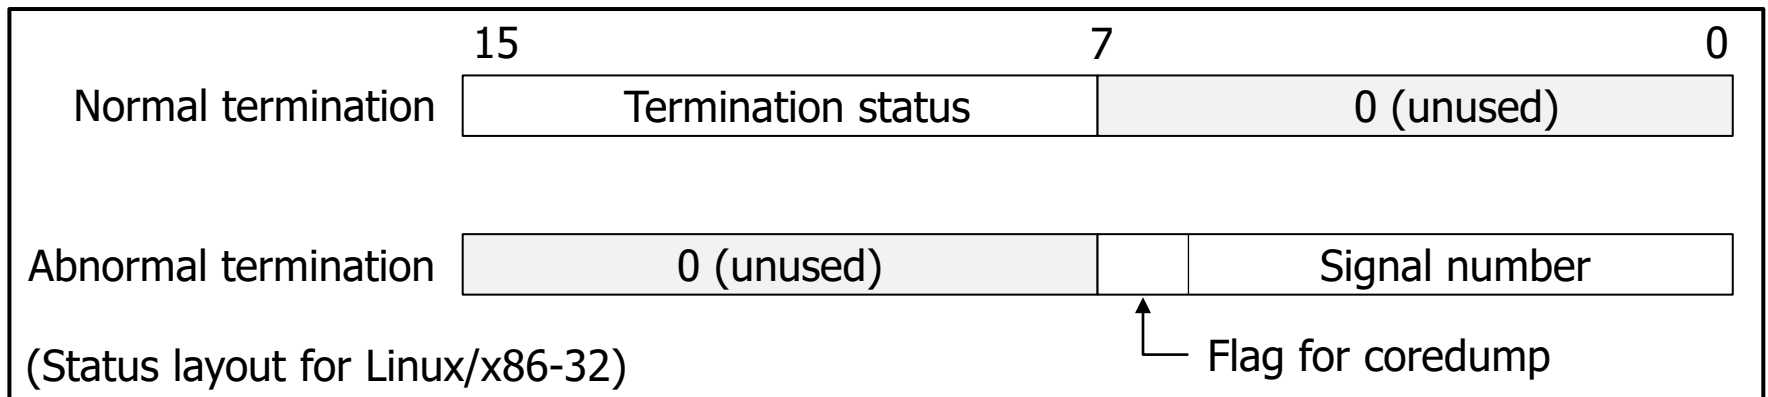
wait (1)

- `#include <sys/types.h>`
`#include <sys/wait.h>`
`pid_t wait(int *status)`
 - Suspends execution of the current process until a child has exited, or until a signal is delivered whose action is to terminate the current process.
 - If the caller blocks and has multiple children, wait returns when one terminates.
 - We can always tell which child terminated because the process ID is returned by the function.

wait (2)

– *status*

- If NULL, wait() ignores the status
- If not NULL, wait() stores the status information (termination status) in the location pointed to by status.
- Only the 16 low order bits are written.
- If wait returns because a child process terminates, the low order 8 bits are 0, the high order 8 bits contains the exit value of the child process.
- If the child exited because of a signal, the high order 8 bits are 0, and the low order 8 bits contain the signal number.



Macros for wait (1)

- WIFEXITED(*status*)
 - Returns true if the child exited normally.
- WEXITSTATUS(*status*)
 - Evaluates to *the least significant eight bits* of the return code of the child which terminated, which may have been set as the argument to a call to `exit()` or as the argument for a return.
 - This macro can only be evaluated if WIFEXITED returned non-zero.

Macros for wait (2)

- WIFSIGNALED(*status*)
 - Returns true if the child process exited *because of a signal*
- WTERMSIG(*status*)
 - Returns *the signal number* that caused the child process to terminate.
 - This macro can only be evaluated if WIFSIGNALED returned non-zero.

Example #3: wait

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>
```

```
int main(void) {
    pid_t pid;
    int status;

    if ((pid = fork()) == 0) { /* child process */
        printf("I am a child\n");
        exit(123);
    }

    /* parent process */
    pid = wait(&status);
    printf("parent: child(pid = %d) is terminated with status (%d)\n", pid,
WEXITSTATUS(status));
    return 0;
}
```

```
❯ ./main
I am a child
parent: child(pid = 133) is terminated with status (123)
```

waitpid (1)

- `#include <sys/types.h>`
`#include <sys/wait.h>`
`pid_t waitpid(pid_t pid, int *status, int options);`
 - waitpid wait for a specific child process, while wait wait for any child process
 - `pid < -1`
 - Wait for any child process whose process group ID is equal to the absolute value of `pid`.
 - `pid == -1`
 - Wait for any child process.
 - Same behavior which `wait()` exhibits.
 - `pid == 0`
 - Wait for any child process whose process group ID is equal to that of the calling process.

waitpid (2)

- *pid* > 0
 - Wait for the child whose process ID is equal to the value of *pid*.
- *options*
 - Zero or more of the following constants can be "OR"ed.
 - WNOHANG
 - » Return immediately if no child has exited.
 - WUNTRACED
 - » Also return for children which are stopped, and whose status has not been reported (because of signal).
 - WCONTINUED
 - » Also return for children which are continued because of SIGCONT.
- Return value
 - The process ID of the child which exited.
 - 0 or ECHILD if WNOHANG was used and no child was available.
 - -1 on error

Macros for waitpid

- WIFSTOPPED(*status*)
 - Returns true if the child process which caused the return is *currently stopped*.
 - This is only possible if the call was done using WUNTRACED.
- WSTOPSIG(*status*)
 - Returns *the signal number* which caused the child to stop.
 - This macro can only be evaluated if WIFSTOPPED returned non-zero.
- WIFCONTINUED(*status*)
 - Returns true if the child process is continued which caused by SIGCONT (Linux 2.6.10 or higher).

Example #4: waitpid (1)

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main() {
    pid_t pid;
    int status, exit_status;

    if ((pid = fork()) < 0)
        perror("fork failed");

    if (pid == 0) {
        printf ("Child %d sleeping... \n", getpid());
        sleep (4);
        exit (5) ;
    }
}
```

Example #4: waitpid (2)

```
while (waitpid(pid, &status, WNOHANG) == 0) {  
    printf ("Still waiting... \n");  
    sleep (1);  
}  
  
if (WIFEXITED (status)) {  
    exit_status = WEXITSTATUS (status);  
    printf ("Exit status from %d was %d\n", pid, exit_status);  
}  
  
exit (0);  
}
```

```
Child 413 sleeping...  
Still waiting...  
Still waiting...  
Still waiting...  
Still waiting...  
Exit status from 413 was 5
```

```
Still waiting...  
Child 226 sleeping...  
Still waiting...  
Still waiting...  
Still waiting...  
Exit status from 226 was 5
```


Exec Functions (1)

- `#include <unistd.h>`
`extern char **environ;`

`int execl(const char *pathname, const char *arg, ...);`

`int execv(const char *pathname, char *const argv[]);`

`int execl(const char *pathname, const char *arg , ..., char *
const envp[]);`

`int execve (const char *pathname, char *const argv[], char
*const envp[]);`

`int execlp(const char *filename, const char *arg, ...);`

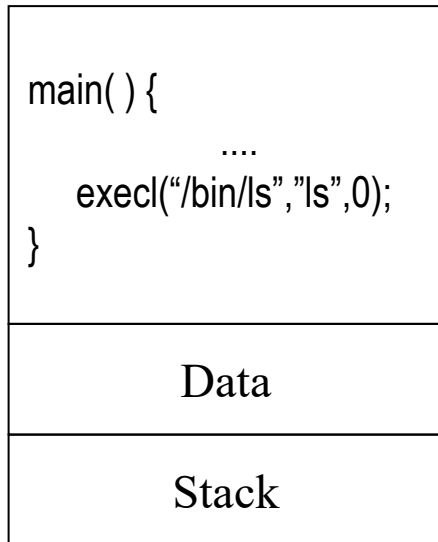
`int execvp(const char *filename, char *const argv[]);`

exec Functions (2)


- Execute program
 - When a process calls one of the exec functions, that process is completely replaced by the new program.
 - The new program starts executing at its `main()` function.
 - `exec` merely replaces the current process (text, data, heap, stack) with a brand-new program from disk.
- Only `execve()` is a system call.
 - `execl`, `execv`, `execle`, `execlp`, `execvp` are front-ends for the function `execve()`.

exec Functions (3)


Before exec pid = 1607



After exec pid = 1607



exec Functions (4)



exec Functions (5)

- Parameters of exec function

- *pathname(filename)*

- The initial argument for these functions is the *pathname (filename)* of a file which is to be executed.
 - `execl`, `execv`, `execle`, `execve` take a *pathname* argument.
 - `execlp`, `execvp` take a *filename* argument.
 - If *filename* contains a slash, it is taken as a *pathname*.
 - Otherwise, the executable file is searched for in the directories specified by the `PATH` environment variable.
 - If `PATH` isn't specified, the default path `“:/bin:/usr/bin”` is used.
 - If permission is denied for a file (the attempted `execve()` returned `EACCES`), these functions will continue searching the rest of the search path.

```
runner@repl.it:~$ echo $PATH
/home/runner/.apt/usr/bin:/usr/local/sbin:/usr/local/bin:/u
sr/sbin:/usr/bin:/sbin:/bin
```

exec Functions (6)

– Argument list

- The const char **arg* and subsequent ellipses in the `execl`, `execlp`, and `execle` functions can be thought of as *arg0*, *arg1*, ..., *argn*.
- Together they describe a list of one or more pointers to null-terminated strings that represent the argument list available to the executed program.
- The first argument, by convention, should point to the file name associated with the file being executed.
- The list of arguments must be terminated by a NULL pointer.

– Example

```
execl("/bin/lis", "lis", "-al", 0);
```

exec Functions (7)

– Argument vector

- `execv`, `execve` and `execvp` provide an array of pointers to null-terminated strings that represent the argument list available to the new program.
- The first argument, by convention, should point to the file name associated with the file being executed.
- The array of pointers must be terminated by a NULL pointer.

– Example

```
char *argv[3] = {"ls", "-al", 0};  
execv("/bin/ls", argv);
```

exec Functions (8)

– Environment variable

- `execve` and `execle` specifies the environment of the executed process by following the NULL pointer that terminates the list of arguments in the parameter list or the pointer to the *argv* array with an additional parameter.
- This additional parameter is an array of pointers to null-terminated strings and must be terminated by a NULL pointer.
- The other functions take the environment for the new process image from the external variable `environ` in the current process.

– Example

```
char *env[2] = {"TERM=vt100", 0};  
execle("/bin/ls", "ls", 0, env);
```


exec Functions (9)

– Environment variable example

```
runner@repl.it:~$ printenv
LC_ALL=en_US.UTF-8
LD_LIBRARY_PATH=/home/runner/.apt/usr/lib/x86_64-linux-gnu:/home/runner/.apt/usr/lib/i386-linux-g
XDG_CONFIG_HOME=/config
LANG=en_US.UTF-8
DISPLAY=MAGIC
HOSTNAME=0320986151f8
OLDPWD=/home/runner/yeonseub
PWD=/home/runner
HOME=/home/runner
CPATH=
LIBRARY_PATH=/home/runner/.apt/usr/lib/x86_64-linux-gnu:/home/runner/.apt/usr/lib/i386-linux-gnu:
APT_OPTIONS=-o debug::nolocking=true -o dir::cache=/tmp/apt/cache -o dir::state=/tmp/apt/state -o
TERM=xterm-256color
SHLVL=1
CPPPATH=
PATH=/home/runner/.apt/usr/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
PS1=\[\033[01;32m\]\u@repl.it\[\033[00m\]:\[\033[01;34m\]\w\[\033[00m\]\$
PKG_CONFIG_PATH=/home/runner/.apt/usr/lib/x86_64-linux-gnu/pkgconfig:/home/runner/.apt/usr/lib/i386-linux-gnu/pkgconfig:/
LD_PRELOAD=/usr/local/lib/repl.so
INCLUDE_PATH=/home/runner/.apt/usr/include:/home/runner/.apt/usr/include/x86_64-linux-gnu:
_=/usr/bin/printenv
```

exec Functions (10)

- Properties that the new program inherits.
 - Open files (depends on *close-on-exec* flag for each file descriptor)
 - pid, ppid, real-user(group)-id
 - Process group id, session id, controlling terminal
 - Current working directory, current root directory
 - Environment variables (except for `execle()`, `execve()`)
- Properties that the new program does not inherits.
 - Signals pending on the parent process are *cleared*.
 - Any signals set by the calling process are *reset* to their default behavior.
- Return value
 - On success, never returns. On error, -1 is returned.

Example #5: execlp

```
#include <sys/types.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdio.h>
#include <sys/wait.h>
```

```
int main(void) {
    pid_t pid;
    int status;

    if ((pid = fork()) == 0) { /* child */
        execlp("ls", "ls", "-a", NULL);
    } else {
        wait(&status);
        printf("exit status of child = %d\n", WEXITSTATUS(status));
    }
}
```

```
➤ ./main
.   bar2   dir   foo   foo-sh  main.c  prog
..  bar-hl  file  foo-d  main    main-sh test
exit status of child = 0
```

Thank you